




# Sequence-to-sequence models with attention



Many slides adapted from [J. Johnson](#)

# Outline

---


- Vanilla seq2seq with RNNs
- Seq2seq with RNNs and attention
- Image captioning with attention

## Sequence-to-sequence modeling: Machine translation


“We are eating bread” → “Estamos comiendo pan”

# Sequence-to-sequence modeling with RNNs

---




I. Sutskever, O. Vinyals, Q. Le, [Sequence to Sequence Learning with Neural Networks](#), NeurIPS 2014



K. Cho, B. Merrienboer, C. Gulcehre, F. Bougares, H. Schwenk, and Y. Bengio, [Learning phrase representations using RNN encoder-decoder for statistical machine translation](#), ACL 2014

# Sequence-to-sequence modeling with RNNs



# Sequence-to-sequence modeling with RNNs


---

**Encoder:**  $h_t = f_W(x_t, h_{t-1})$


From final hidden state predict:

**Initial decoder state**  $s_0$


**Context vector**  $c$  (often  $c = h_T$ )



# Sequence-to-sequence modeling with RNNs



# Sequence-to-sequence modeling with RNNs






# Sequence-to-sequence with RNNs *and attention*

---


- Intuition: translation requires *alignment*




# Sequence-to-sequence with RNNs and attention

---


- At each timestep of decoder, context vector “looks at” different parts of the input sequence




# Sequence-to-sequence with RNNs and attention




# Sequence-to-sequence with RNNs and attention



# Sequence-to-sequence with RNNs and attention

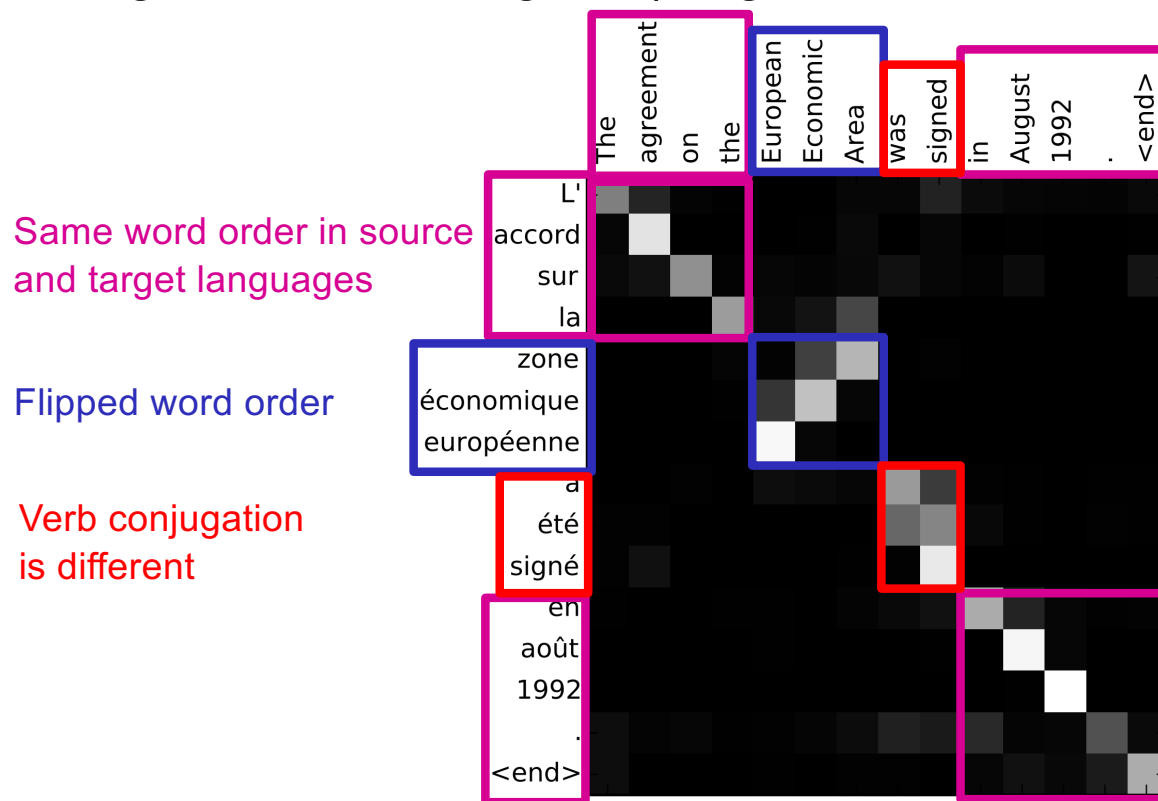


# Sequence-to-sequence with RNNs and attention




# Sequence-to-sequence with RNNs and attention

- Visualizing attention weights (English source, French target):



# Quantitative evaluation

---





# Google Neural Machine Translation (GNMT)

---

## Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation


Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le, Mohammad Norouzi  
`yonghui,schuster,zhifengc,qvl,mnorouzi@google.com`

Wolfgang Macherey, Maxim Krikun, Yuan Cao, Qin Gao, Klaus Macherey,  
Jeff Klingner, Apurva Shah, Melvin Johnson, Xiaobing Liu, Łukasz Kaiser,  
Stephan Gouws, Yoshikiyo Kato, Taku Kudo, Hideto Kazawa, Keith Stevens,  
George Kurian, Nishant Patil, Wei Wang, Cliff Young, Jason Smith, Jason Riesa,  
Alex Rudnick, Oriol Vinyals, Greg Corrado, Macduff Hughes, Jeffrey Dean

Y. Wu et al., [Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation](#), arXiv 2016

<https://www.nytimes.com/2016/12/14/magazine/the-great-ai-awakening.html>

# Google Neural Machine Translation (GNMT)



Y. Wu et al., [Google's Neural Machine Translation System: Bridging the Gap between Human and Machine Translation](#), arXiv 2016

# Google Neural Machine Translation (GNMT)

---

- **Standard training objective:** maximize log-likelihood of ground truth output given input:

$$\sum_i \log P_W(Y_i^* | X_i)$$

- Only encourages the system to reproduce the reference sentences, does not induce a very good ranking on outputs that don't match reference sentences
- Not related to task-specific evaluation metric (e.g., BLEU score)
- **Refinement objective:** expectation of rewards over possible predicted sentences  $Y$ :

$$\sum_i \sum_Y P_W(Y | X_i) R(Y, Y_i^*)$$

- Use variant of BLEU score to compute reward
- Reward is not differentiable -- need RL to train (initialize with ML-trained model)

# Google Neural Machine Translation (GNMT)

---

- Human evaluation results on production data (500 randomly sampled sentences from Wikipedia and news websites)

Table 10: Mean of side-by-side scores on production data

|                   | PBMT  | GNMT  | Human | Relative Improvement |
|-------------------|-------|-------|-------|----------------------|
| English → Spanish | 4.885 | 5.428 | 5.550 | 87%                  |
| English → French  | 4.932 | 5.295 | 5.496 | 64%                  |
| English → Chinese | 4.035 | 4.594 | 4.987 | 58%                  |
| Spanish → English | 4.872 | 5.187 | 5.372 | 63%                  |
| French → English  | 5.046 | 5.343 | 5.404 | 83%                  |
| Chinese → English | 3.694 | 4.263 | 4.636 | 60%                  |

**Side-by-side scores:** range from 0 (“completely nonsense translation”) to 6 (“perfect translation”), produced by human raters fluent in both languages

**PBMT:** Translation by phrase-based statistical translation system used by Google

**GNMT:** Translation by GNMT system


**Human:** Translation by humans fluent in both languages

# Outline

---

- Vanilla seq2seq with RNNs
- Seq2seq with RNNs and attention
- Image captioning with attention


# Generalizing attention




# Image captioning with RNNs and attention

---

- Idea: pay attention to different parts of the image when generating different words
- Automatically learn this *grounding* of words to image regions without direct supervision




# Image captioning with RNNs and attention




Use CNN to extract a grid of features




# Image captioning with RNNs and attention



# Image captioning with RNNs and attention



# Image captioning with RNNs and attention



# Example results

---

- Good captions



A woman is throwing a frisbee in a park.



A dog is standing on a hardwood floor.



A stop sign is on a road with a mountain in the background.



A little girl sitting on a bed with a teddy bear.



A group of people sitting on a boat in the water.



A giraffe standing in a forest with trees in the background.


# Example results

---

- Mistakes



A large white bird standing in a forest.



A woman holding a clock in her hand.



A man wearing a hat and  
a hat on a skateboard.



A person is standing on a beach  
with a surfboard.



A woman is sitting at a table  
with a large pizza.



A man is talking on his cell phone  
while another man watches.

# Quantitative results

---

| Dataset   | Model          | BLEU-1      | BLEU-2      | BLEU-3      | BLEU-4      | METEOR       |
|-----------|----------------|-------------|-------------|-------------|-------------|--------------|
| Flickr8k  | Google NIC     | 63          | 41          | 27          | -           | -            |
|           | Soft-Attention | <b>67</b>   | 44.8        | 29.9        | 19.5        | 18.93        |
|           | Hard-Attention | <b>67</b>   | <b>45.7</b> | <b>31.4</b> | <b>21.3</b> | <b>20.30</b> |
| Flickr30k | Google NIC     | 66.3        | 42.3        | 27.7        | 18.3        | -            |
|           | Soft-Attention | 66.7        | 43.4        | 28.8        | 19.1        | <b>18.49</b> |
|           | Hard-Attention | <b>66.9</b> | <b>43.9</b> | <b>29.6</b> | <b>19.9</b> | 18.46        |
| COCO      | Google NIC     | 66.6        | 46.1        | 32.9        | 24.6        | -            |
|           | Soft-Attention | 70.7        | 49.2        | 34.4        | 24.3        | <b>23.90</b> |
|           | Hard-Attention | <b>71.8</b> | <b>50.4</b> | <b>35.7</b> | <b>25.0</b> | 23.04        |

[Source](#)